

Annotation Guidelines

Spotting Stopping Points (i.e. Critical Interventions) in Online Conversations

This document outlines the definition of the notion of a **stopping point** and provides annotation guidelines to determine whether a user comment on a social media constitutes a stopping point or not.

TABLE OF CONTENT

1. Definition of stopping point and unit of analysis.....	2
1.1. Definition.....	2
1.2. Unit of analysis.....	2
2. Sampling procedure.....	4
3. Annotation procedure and principles.....	4
3.1. Annotation procedure.....	4
3.2. Annotation principles.....	5
4. Complicated or ambiguous cases — resolution rules and examples.....	5
Case 1 — Refutation or reframing without explicit lexical markers.....	5
Case 2 — Incredulous alignment.....	6
Case 3 — Humor, parody or irony.....	6
Case 4 — Very short or fragmentary comments.....	7
Case 5 — Parent post already performs the refutation.....	7
Case 6 — Parent post is reported speech.....	8
Case 7 — Replies to a stopping point comment.....	9
Case 8 — Same user posts several comments.....	9
Additional complicated cases to consider.....	10
5. Test for annotators.....	10

1. Definition of stopping point and unit of analysis

1.1. Definition

A **stopping point** is a user's critical intervention in an online discussion¹. It covers a range of forms — from a brief expression of doubt (“Is this true?”) or a dismissive remark (“You’re talking nonsense”) to a more elaborate refutation or correction with counter-arguments and/or sources. The act itself does not imply the factual accuracy or normative legitimacy of the intervention. Tone, rhetoric or style do not disqualify an utterance: angry, sarcastic, insulting or informal language can still count as a stopping point if it functions as a critical intervention in the thread.

Stopping points may target different aspects of a post or discussion, such as:

- **Credibility of content:** *This is completely false! / fake news*
- **Relevance or interest:** *This is nonsense / bullshit / who cares*
- **Source reliability:** *This source is completely biased / Is there still any real journalism?*
- **Form of content** (e.g., images, videos): *Fake Photoshop / badly done montage*
- **The poster:** *Your page is shameful / I'm unfollowing*
- **Other users:** *Can't believe people actually fall for this / Enough with your worthless comments*

1.2. Unit of analysis

Stopping point may appear either as a **comment** or as a **nested reply** to another comment. The **unit of analysis** is thus the **single textual contribution** (comment or reply) considered in its local interactional context.



Figure 1. Examples of stopping points appearing either as a comment (left screenshot) or as a nested reply to another comment (right screenshot)

¹ Although the concept of a stopping point could also be applied to offline conversations, it has so far only been operationalized for the analysis of online interactions. For this reason, we define it here specifically in relation to digital conversational contexts.

Importantly, a stopping point is defined as a **momentary discursive act**, not as a stable position of the author. One user may later repeat a stopping point to strengthen its position, or may retract, nuance, or reverse its stance. Conversely, one user who initially posts an aligned or neutral comment can subsequently produce a stopping point. Stopping points should thus be treated as a **conversational pause** — a brief interruption or timeout that can trigger rebuttals, new arguments, jokes, or further clarification — rather than as a **definitive verdict that necessarily closes the discussion** (even though in some threads it may have that effect).



Figure 2. Same user expresses two stopping points: first in a comment, then in a nested reply

Translation and explanations:

- Top-level comment: “Did you read the article before posting it? Obviously not — you would have quickly seen it’s a hoax.” → **YES**: Explicitly labels the post as a hoax and questions whether the author checked the source.
- Other reply: “You’re telling me! ripping ticks out with a crowbar!” → **NO**: purely humorous.
- User’s follow-up reply: “lol or a tick ate an entire hand. Still, it’s serious to publish without bothering to read.” → **YES**: Although phrased with irony and laughter, the reply performs a procedural critique (denouncing negligent publication)



Figure 3. Same user expresses one stopping point in a comment, then a conversational nested reply

Translation and explanations:

- Top-level comment: “There are still some bullshit things in what they say.” → **YES**: direct dismissal of content relevance/credibility (targets the post’s substance).
- Reply: “Well, it’s Topito after all...” → **NO**: conversational/justificatory remark about the source’s character; functions as mitigation/affiliation rather than an attempt to halt, verify or correct.
- User’s follow-up reply: “Yeah, except afterwards I end up with people who give me their PC and then doubt my skills because ‘oh but I read on Topito that it doesn’t work like that’.” → **NO**: conversational anecdote/complaint; it elaborates a personal consequence rather than directly contesting the post’s truth or calling for verification/moderation.



Figure 4. Same user expresses first a credulous comment, then a stopping point in a nested reply

Translation and explanations:

- **Top-level comment:** “What I don’t understand is that nobody noticed this in four years and that above all this lunatic gets only a fine, 2 years in prison that he may not even serve entirely — and above all nobody talks about compulsory treatment even though this guy is really sick...” → **NO**: this is a credulous, outraged reaction that treats the content as factual and calls for sanction; it does not contest the post’s truth or request verification
- **Reply:** “No doubt because it’s a parody news site... ‘BMF News’ not ‘BFM’... so Fake News.” → **YES**: explicit source-based refutation: identifies parody / incorrect source and labels the content as fake.
- **User’s follow-up reply:** “Ok, thanks. That kind of post is really rubbish.” → **YES**: explicit evaluative critique of the post (relevance / quality). Even though short and informal, it functions as a procedural denouncement and thus counts as a stopping point.

2. Sampling procedure

Before starting the annotation, select a random sample of comments representing a fixed proportion of the total corpus (e.g., 1% to 10%). We recommend randomly selecting posts and annotating all comments associated with each post, rather than sampling isolated comments. This approach makes it possible to follow complete discussion threads directly on the platform and to interpret each comment within its full interactional context.

3. Annotation procedure and principles

3.1. Annotation procedure

For each comment:

1. **Read the parent post** (title/content) and **check the media source** and the **page/group description** to understand context and tone.
2. **Read adjacent comments** (preceding and following comments/replies) when available.
3. Apply the decision rule:
 - Label **YES** if the comment explicitly expresses criticism, correction, or a call for verification/action aimed at any of the items below :
 - **Credibility of content:** “*this is completely false! / fake news*”
 - **Relevance or interest:** “*This is nonsense / bullshit / who cares.*”
 - **Source reliability:** “*Is there still any real journalism?/ this source is completely biased*”
 - **Form of content** (e.g., images, videos): “*Fake Photoshop, badly done montage.*”
 - **The poster:** “*Your page is shameful, I’m unfollowing.*”
 - **Other users:** “*Can’t believe people actually fall for this.*”

- Label **NO** if the comment is neutral, supportive, purely conversational, humorous without critical intent, or off-topic. Comments that only tag other users without any other words can also be quickly classified as **NOT** a stopping point.
4. Record brief notes for ambiguous cases (if a notes field is available).

3.2. Annotation principles

Annotators are tasked with identifying units of text that constitute stopping points, without evaluating truth, validity, sincerity, motivation, or normative quality. Their role is solely to detect whether a text expresses a critical stance or calls for action/verification regarding the post, its source, presentation, the poster, or other users. In practice, any explicit or implicit attempt to halt, question, correct, denounce, or prompt verification/moderation should be labeled as a stopping point, regardless of register — whether formal argument, insult, emoji, or rhetorical question. Accuracy or constructive quality is irrelevant: even comments that rely on inaccurate claims or use harmful discourses qualify. These interventions, though potentially unconstructive, interrupt circulation, elicit responses, and shape the negotiation of norms and collective meaning in the thread — the very dynamics that the concept of stopping point aims to capture.

4. Complicated or ambiguous cases — resolution rules and examples

This non-exhaustive list presents cases that annotators may find challenging: subtle critiques that particularly require careful, context-sensitive reading, and ambiguous interventions that raise genuine doubts or questions for the annotator. Each case below gives (a) a short description of the difficulty, (b) contextual cues or indices to rely on before making a decision, (c) the resolution rule to apply, and (d) subtle examples and annotations with brief justifications. Research teams should run calibration exercises on these cases and may adapt the rules to their objectives.

Case 1 — Refutation or reframing without explicit lexical markers

- **Description:** The comment implies rejection, opposition, correction, or reframing but lacks clear lexical cues (e.g. no “fake”, “source”, “hoax”, etc.).
- **Cues:** The comment’s statement contradicts the post ; diverges from the publisher’s stance ; contains warning/consequence language, or uses an imperative/discouraging tone — implying that if the post were true it would lead to negative outcomes, thus signaling the commenter’s critique.
- **Resolution rule:**
 - **YES** if the contribution offers a counter-argument, expresses an opposing opinion, reframes the claim, or warns about its consequences.
 - **NO** if it only adds related information that does not call into question the original claim.
- **Examples:**
 - Post: “Don’t vaccinate your kids!” ; Comment: “Don’t come crying later!”
→ **YES**: the comment implicitly signals that the recommendation is dangerous or irresponsible and forecasts negative outcomes if followed. That forecasting functions as a corrective move: it discourages the action and raises doubt about its safety/validity.
 - Post: “Walking reduces cardiovascular risk by 15%.” ; Comment: “The effect is stronger in cities”
→ **NO**: the comment provides additional information that does not contradict the original claim.

Case 2 — Incredulous alignment

- **Description:** The comment aligns with the post's stance but expresses conditional doubt (e.g., "if this is true"), often with outrage. This type of reaction often **does not count** as stopping points but can be described as "incredulous alignment" for fine-grained annotation since the critical gesture is subordinated to alignment or outrage rather than a clear attempt to question, halt, or correct the content.

- **Contextual cues:** presence of "if true/if that's the case" or of conjunctions such as "but" or "still" — often treated in linguistic as an interactive marker situated in an intermediate zone between agreement and refutation (Moeschler & De Spengler, 1981); heavy affective language; calls for punitive action without evidence or without requesting sources.

- **Resolution rule:**

- **NO** if the main function is conditional outrage/alignment.

- **YES** only if it explicitly requests verification, provides sources, or calls for moderation.

- **Examples:**

- "If this is true, he should be removed from office!"

→ **NO:** outrage only.

- "If true, that's outrageous — can anyone cite the original speech?"

→ **YES:** verification request.

Case 3 — Humor, parody or irony

- **Description:** When the parent post or its surrounding context is explicitly satirical or humorous, the publication itself often functions as a provocation: disagreement, mockery, or derision are part of the expected reception and, much like in a debate, may be perceived as "preferred responses" (Kerbrat-Orecchioni, 2016). Crucially, however, not all forms of critique are "preferred": some do constitute genuine stopping points. Irony and second-degree humor therefore occupy a grey zone: they may serve purely social functions (e.g. signaling amusement, group belonging, or ridicule) or perform genuine critical work (denouncing falsity, prompting verification, or distancing from a viewpoint - see Baklouti, 2016 for a linguistic analysis of irony as disagreement). Annotators must therefore distinguish between playful critique (affiliative, social, or purely derisive) and operative ironic moves that functionally interrupt, question, or correct the post.

- **Resolution rule.**

→ **YES:** If irony explicitly functions to denote incredulity and prompts verification or correction

→ **NO:** If irony is purely affiliative, humorous, or marks suspension of disbelief

- **Examples:**

- "LOL 😂" / "Haha I can't take it anymore"

→ **NO (phatic humor):** comments may signal non-credulity but that disbelief is expressed as play or amusement rather than as a corrective or verification act.

- "Damn, I wanted to believe it!" / "Too bad, that would have been hilarious."

→ **NO (suspension of disbelief):** Users recognize that the published content is not factual but satirical but they play along with the idea that it *could* be true, and sometimes even express the wish that it were. The users momentarily entertain or enjoy the plausibility of the satire but do not seek to verify, correct, or halt its circulation.

- Post: Michel Sardou Died ; Comments: "Damn, he died again", "Michel Sardémarre", "How many times

a year does he die?”

→ **NO (playful parody reuse)**: Commenters are not credulous but disbelief here is enacted as play: sarcasm and puns perform group membership and communal recognition of past hoaxes.

- “Great, more of that”

→ mostly **NO (Implicit sarcasm that mimics endorsement)** : could be ironic or sincere; rely on thread cues and author history if available.

- “fake bro” ; “dude you have nothing better to do than share dumb posts like this? hahaha”

→ **YES (social but still corrective)**

“Yeah, and pigs will fly next week 😂.”

→ **YES (irony as refutation)**: irony denoting incredulity

- “Not funny in times like these.”

→ **YES (explicit rejection of humor)**: contests the legitimacy of the post itself.

Case 4 — Very short or fragmentary comments

- **Description**: Many online comments consist of just a single word or fragment. Some of these are unambiguous stopping points, such as “Fake” or “Source?”, which directly contest truth or request verification. Others are more ambiguous: link-only comment, mere emoji denoting doubts like 🤔 or 🙄, minimal interjections like “pfff”, a bare date like “2008”, or a single mention of a source (“Le Gorafi”) without further explanation.

- **Contextual cues**: In such cases, interpretation depends heavily on context. Longer surrounding comments can clarify the meaning and other users expand the fragment into a clearer critique.

- **Resolution rules**:

- **YES** if surrounding context indicates the fragment functions as contestation or a source correction

- **NO** if no contextual cues exist

- **Examples**:

- If other users write “Why repost an article from 2008?”, then a standalone “2008” likely serves as a corrective marker.

- If others warn “Look carefully, it’s BFM and not BFM TV,” then “bfm” alone can be read as pointing out source reliability.

- If other comments denounce the post with “pff this is nonsense, don’t believe everything you see on social media,” then a minimal “pff” in the same thread may reasonably be interpreted as implicit contestation. But Isolated “pff” with no clarifying context → **NO** (phatic interjection).

- Link-only comment or reply : if labeled as a fact-check or clearly contradict precedent claim **YES** otherwise **NO**

- If other comments contest the post, then emoji like 🤔 or 🙄 can reasonably be interpreted as doubts or critics

Case 5 — Parent post already performs the refutation

- **Description**: Parent post already corrects or debunks the content; comments may either add to that correction or merely echo it.

- **Contextual cues**: Does the reply introduce new evidence, a link, a date, or a procedural action (e.g., “I reported this”)? Or is it just an affirmation/endorsement?

- **Resolution rule**:

→ **YES** if reply adds new corrective content (e.g., link, evidence, “I reported this”).

→ **NO** if reply simply agrees (“Exactly!”, “+1”).

- **Examples:**

- **Post:** “False — here’s a debunk.” ; **Comment:** “[source attached]”

→ **YES:** because it adds evidence.

- **Post:** “False.” ; **Comment:** “Exactly!”

→ **NO:** because simple agreement.

Case 6 — Parent post is reported speech

- **Description:** The parent post republishes or quotes an external claim (e.g., “According to X, Y happened,” or a shared screenshot of a quote). In such cases, comment may react to different targets: the *quoted claim itself*, the *source being quoted*, or the *act of reposting/republishing* (tone, framing, selection). Correct annotation requires identifying which target the comment addresses.

- **Contextual cues:** Look for lexical cues (pronouns, source names, quotation marks, verbs like “they say”, “quote”, “reported”) to identify if the comment target the media, the quoted claim or the source being quoted, the page/group who repost the content, etc. Check the editorial line of the media and/or the description of the page/group, as well as recent posts to identify whether they are aligned or not with the quoted claim. Look also at surrounding turns: do adjacent comments clarify whether the thread debates the quoted claim, the source, or the reposting action?

- **Resolution rule:**

- If the media’s editorial line or the page description makes clear that the repost *endorses* the quoted claimed content and if the reply criticises the quoted claim or the act of sharing (e.g. “You shouldn’t repost this”), label **YES** (the reply functions to contest and potentially halt spread).

- If the media’s editorial line or the page description makes clear that the repost is already critical or satirical, then a reply that simply reiterates that critique (i.e. also criticizes the quoted claim) should be marked **NO**. Conversely, if the reply criticises the *act of sharing* (e.g. “don’t repost this”) or attacks/denigrates the person who reposted it in a way that aims to curb circulation or call out the re-poster, mark **YES**.

- **Examples:**

- **Post:** “Professor XX says” ; **Context:** the media’s editorial line or the page description make clear that the repost endorses the quoted claim ; **Comment:** “This claim is misleading! Professor X was struck off the medical register.”

→ **YES:** because the comment directly contests the claim and supplies corrective evidence

- **Post:** “X says immigrants are taking all jobs.” ; **Context:** the media’s editorial line or the page description make clear that the repost is already critical or satirical towards the quoted claim ; **Comment:** “Aren’t you ashamed to give publicity to far-right propaganda?”

→ **YES** because the comment criticizes the **act of sharing** and calls out the re-poster

- **Post:** “X says immigrants are responsible for spreading viruses.” ; **Context:** the repost comes from a clearly satirical account mocking the claim ; **Comment:** “He’s so dishonest he’d be capable of saying something like that.”

→ **NO** because the reply merely echoes or reinforces the page’s critical/satirical framing

Case 7 — Replies to a stopping point comment

- **Description:** A potential stopping point appears as a reply to another comment (not to the original post). The reply may contest the original post indirectly or the parent comment.
- **Contextual cues:** identify which antecedent the reply targets (the post vs. the commenter), check whether the reply **adds evidence** (link, citation, screenshot), **asks for verification**, or **announces an action** (e.g. “I reported this”) ; consult adjacent turns: do surrounding replies clarify whether the reply is a stopping point or not.
- **Resolution rule:**
 - **YES** if reply adds new corrective content, e.g. supplies a link or source, provides factual evidence, requests verification, calls for reporting/removal, or otherwise advances the epistemic status of the claim.
 - **NO** if the reply simply agrees with the parent comment (“Exactly!”, “+1”).
 - Mainly **NO** if the reply contests the initial stopping point *without* bringing new corrective information. **Exception:** if the defence includes new arguments or sources, treat as **YES**.

Examples:

- **Parent comment (A):** “Stop with your anti-vax propaganda” ; **Reply (B):** “Anti-vax sweepers.”
 - **NO:** B simply echoes or amplifies A’s stance without providing evidence, correction, or a request to verify; it rallies or mocks rather than functioning as a corrective intervention.
- **Parent comment (A):** “Stop with your anti-vax propaganda” ; **Reply (B):** “This is also dangerous. Measles outbreaks have resurged in some countries. See this report: [link]”
 - **YES:** B supplements the thread with argument and an external source; it adds corrective content and aims to inform/verify.
- **Parent comment (A):** “This is false” ; **Reply (B):** “No — this is satirical”
 - **NO:** B reiterates or reinforces the page’s ironic/satirical framing but does not provide corrective evidence or call for verification;
- **Parent comment (A):** “This is false”; **Reply (B):** “Yeah and this isn’t funny at all; it fuels hateful narratives. I’ve reported it to the moderators.”
 - **YES:** Reply B combines a critique of the content (highlights its harm) with a procedural action (reporting), thus aiming to interrupt circulation and trigger verification.

Case 8 — Same user posts several comments

- **Description:** A single user may post several successive comments or replies in the same thread. These can switch function over time (e.g., from a neutral/phatic remark → to a critical intervention, or from a stopping point → to a retraction/nuance). Because a stopping point is a momentary discursive act, the order and content of each turn matter for annotation.
- **Contextual cues:** timestamps/order, explicit markers of change of mind (e.g., “actually”, “sorry”, “I was wrong”, “no wait”), whether the later comment clearly addresses veracity, supplies evidence or requests moderation.
- **Resolution rules:**
 - **YES (1)** only the message(s) whose primary function is to interrupt, question, correct, request verification, or call for procedural action.
 - **NO (0)** the user’s turns that are neutral, phatic, supportive, or that retract/undermine their earlier critical act.

- If a user posts an initial stopping point and later extends it with evidence, both turns may be coded YES (first as the initial intervention, second as supporting corrective material). If the user retracts the stopping point, the retraction is NO.

- **Examples:**

- **YES then YES (extension):**

- Comment 1: "This is fake."

- **YES**

- Comment 2 (same user): "See this study: [link] — it disproves the claim."

- **YES** (extends the stopping point with evidence)

- **YES then NO (retraction):**

- Comment 1: "This is fake, see link."

- **YES**

- Comment 2 (same user): "Oh, my mistake — I didn't realize it was satire. Sorry!"

- **NO** (retraction)

- **NO then YES (later correction):**

- Comment 1: "Wow! Super cool"

- **NO** (phatic/neutral)

- Comment 2 (same user): "Wait, that clip is actually from 2009 — misleading."

- **YES** (later corrective turn)

Additional complicated cases to consider

- **Quoted text / repetition** — User quotes the post and adds no comment: usually **NO** unless the quote is prefaced by verification language. Could be ironic

- **Cross-lingual / code-switching ambiguity** — Nonstandard language or code-switching can hide markers; use cultural/contextual knowledge and adjacent comments.

5. Test for annotators

Before launching a full annotation campaign and calculating inter-annotator agreement, it is essential to verify that the guidelines are sufficiently clear and operational. To this end, a short pilot test can be conducted on a sample of 10 cases deliberately selected for their ambiguous or non-obvious nature.

The table below reports the annotations produced by each annotator, along with final decision and justification after discussion in case of disagreement, in order to document reasoning processes, identify points of divergence, and refine the guidelines where necessary. Here, there were three instances of disagreement. After discussion annotators 2 and 3 accepted annotator 1's labels, which allowed us to refine the guidelines for comments posted in reply to parodic posts (case 3) or in reply to stopping point (case 7).

Media Source	Parent Post	Comment_type	Text	Annotation 1	Annotation 2	Annotation 3	Final decision in case of disagreement
Les Moutons Rebelles	Another two-month-old baby dies 48 hours after receiving eight vaccines	comment	At the same time, before vaccines thousands of children used to die. So it's not perfect, but oh well!	stop	stop	stop	
Science Info	Alert: carnivorous Asian ticks reported in several French departments	reply	@XXX as far as I know that's humour. Our friend @XXX loves to bring good vibes to this group. So she shares anything that's funny	no_stop	no_stop	no_stop	
BMF News	Food service: He ejaculated into the tacos before serving them to customers!	reply	BMF	stop	stop	stop	
BMF News	Food service: He ejaculated into the tacos before serving them to customers!	reply	Fake, but funny	stop	stop	stop	
Prévention Santé	Vaccinated children more affected by chronic diseases than unvaccinated children	reply	Anti-vax sweepers.	no_stop	stop	stop	no_stop : the reply echoes or amplifies previous stopping points without adding argument, source, or verification. It merely rallies or mocks
Nord Presse	If the Yellow Vests worked more instead of complaining online, they would be less poor	comment	Oh! is the 'taking it literally' team already here?	no_stop	stop	stop	no_stop : the comment criticizes other stopping points by siding with the paradoxical post's framing), not by contesting it
Sain et Naturel	Trump has just authorized the slaughter of bears and wolves, including their cubs, in Alaska	comment	I hope that's false!!! 😞😞😞	no_stop	no_stop	no_stop	
Science Info	Senators' Christmas bonus will be capped this year at only €17,580	comment	It's a hoax — no, it's a parody and it's funny! By the way, MPs get €1,200 in housing benefits in Paris and they recently gave themselves a €300 raise! Isn't that factual, Mr. government fact-checkers? Come on, happy holidays! With your meal expenses alone I could feed a table of 100 people at Christmas.	no_stop	stop	stop	no_stop : the comment criticizes other stopping points by reinforcing the post's ironic stance (mocking parliamentarians)
Santé + Mag	Michael Schumacher is no longer confined to his bed, five years after his skiing accident	comment	😞😞😞😞😞	stop	stop	stop	
La Gauche m'a tuer	Maxence — robbed, beaten and molested by 15 migrants. Total media silence smothering the case.	reply	Okay, you didn't understand the discussion at all, so I'll recap: I almost agree with you. What I'm criticizing is the crude manipulation in articles like this. Apart from making everyone who follows this page look like idiots, it's useless. The page was shut down for a month, and instead of falsifying information it should publish real, up-to-date information to be credible.	stop	stop	stop	